

# Linux Performance Tuning

**Thomas Cameron, RHCA, AWS CSAA**



**thomas.cameron@camerontech.com**



**@thomasdcameron**

# Agenda

- **Intro**
- **What is performance tuning?**
- **Areas to tune**
- **Baseline**
- **Tools for tuning**



# Intro

@thomasdcameron | | thomas.cameron@camerontech.com

# Intro

- I'm Thomas!

# Intro

- **I'm Thomas!**
- **Been in IT since '93**

# Intro

- **I'm Thomas!**
- **Been in IT since '93**
- **Been working with Linux since '95**

# Intro

- I'm Thomas!
- Been in IT since '93
- Been working with Linux since '95
- Performance tuning is a science, not an art

# Intro

- I'm Thomas!
- Been in IT since '93
- Been working with Linux since '95
- Performance tuning is a science, not an art
- Anyone who says different doesn't know how to do it right

# What is performance tuning?

# What is performance tuning?

- **Performance tuning is iterative**

# What is performance tuning?

- **Performance tuning is iterative**
- **You make a baseline**



# What is performance tuning?

- **Performance tuning is iterative**
- **You make a baseline**
- **You find the worst bottleneck**

# What is performance tuning?

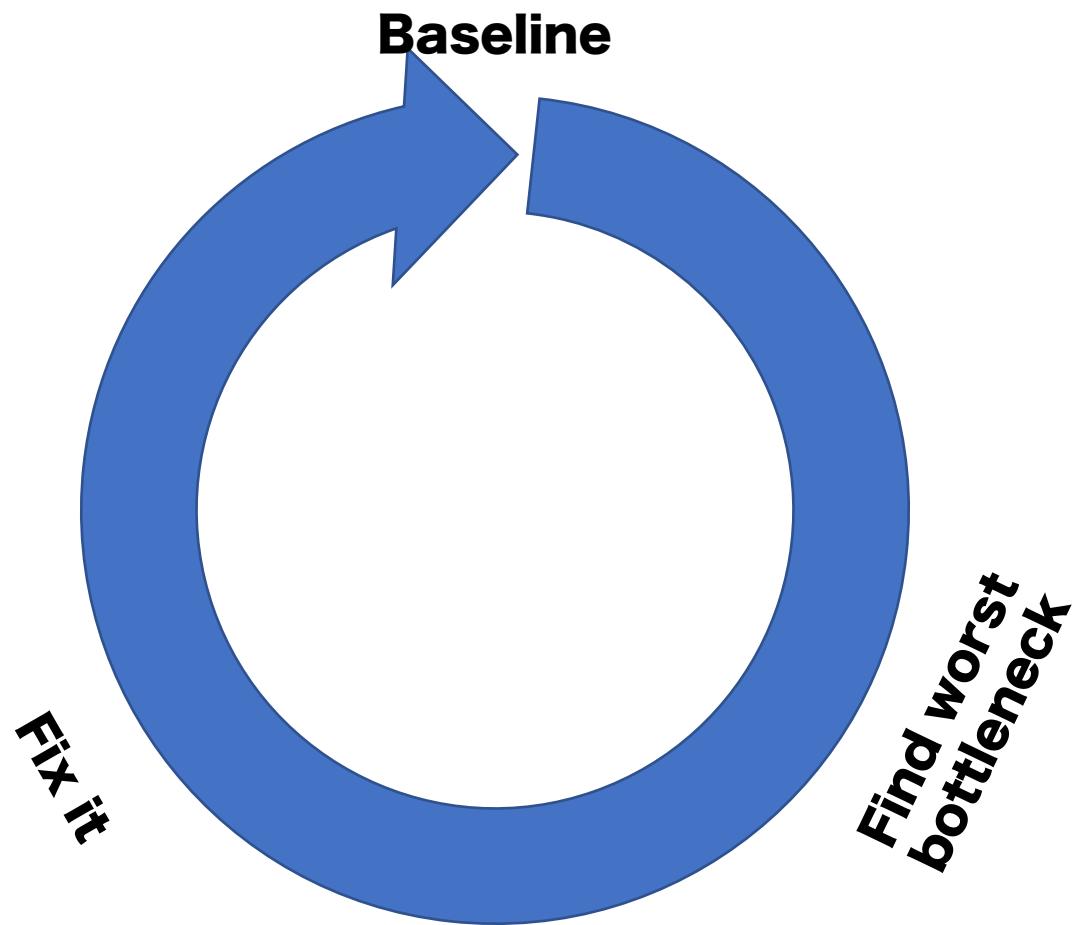
- **Performance tuning is iterative**
- **You make a baseline**
- **You find the worst bottleneck**
- **You fix it**

# What is performance tuning?

- **Performance tuning is iterative**
- **You make a baseline**
- **You find the worst bottleneck**
- **You fix it**
- **You create a new baseline**

# What is performance tuning?

- **Performance tuning is iterative**
- **You make a baseline**
- **You find the worst bottleneck**
- **You fix it**
- **You create a new baseline**
- **Lather, rinse, repeat**





# Areas to tune

# Areas to tune

- **There are 4 subsystems we need to look at**

# Areas to tune

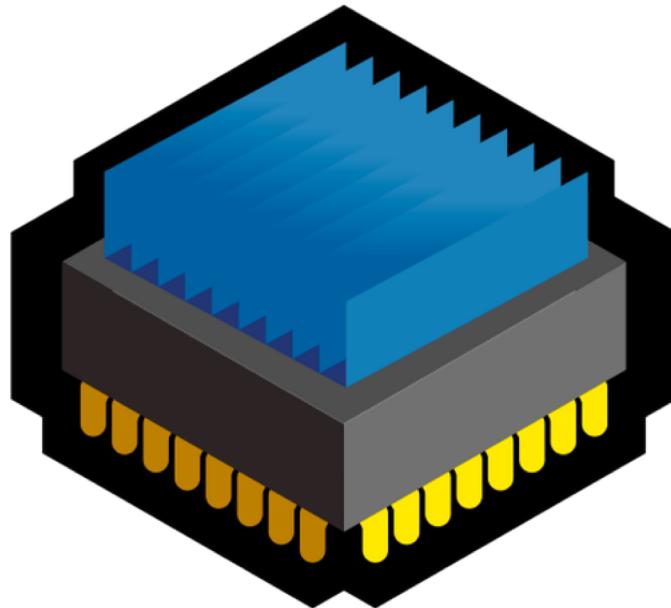
- **There are 4 subsystems we need to look at**
- **Some we can fix through OS level changes**

# Areas to tune

- **There are 4 subsystems we need to look at**
- **Some we can fix through OS level changes**
- **Some are more hardware fixes than anything**

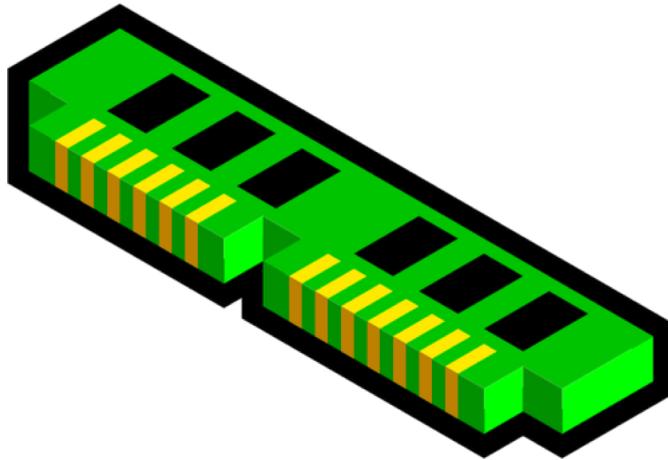
# Areas to tune

- CPU



# Areas to tune

- CPU
- Memory



# Areas to tune

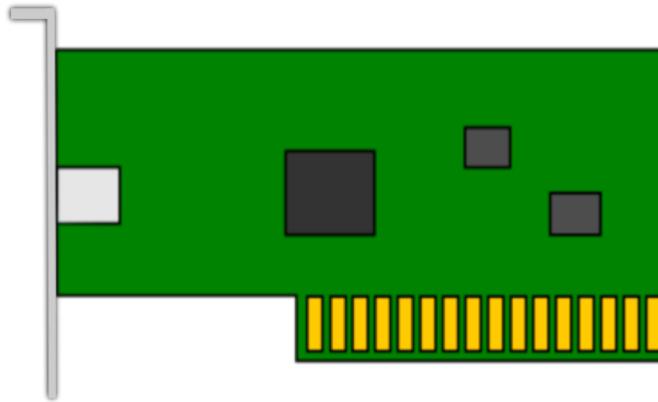
- CPU
- Memory
- Disk





# Areas to tune

- **CPU**
- **Memory**
- **Disk**
- **Network**





# Make a baseline

# Make a baseline

- **Run the app or service normally**

# Make a baseline

- **Run the app or service normally**
- **Gather metrics on how your system is performing**



# Tools for baseline

# Tools for baseline

- There are a number of tools you can use

# Tools for baseline

- **There are a number of tools you can use**
- **We'll go over a few of them**

# Tools for baseline

- **ps**

# Tools for baseline

- **ps**
  - **Don't underestimate this little utility!**

# Tools for baseline

- **ps**
  - **Don't underestimate this little utility!**
  - **Use this to see how many and what processes are running**

A screenshot of a terminal window titled 'Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...'. The window shows a root shell on an AWS EC2 instance. The user runs the command 'ps' to list processes. The output is as follows:  
[root@ip-172-31-45-240 ~]# ps  
 PID TTY TIME CMD  
29525 pts/0 00:00:00 sudo  
29526 pts/0 00:00:00 bash  
29545 pts/0 00:00:00 ps  
[root@ip-172-31-45-240 ~]#  
The terminal has a standard OS X-style interface with red, yellow, and green window control buttons at the top left. A vertical scroll bar is visible on the right side of the terminal window.



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...  
[root@ip-172-31-45-240 ~]# ps  
  PID TTY      TIME CMD  
29525 pts/0    00:00:00 sudo  
29526 pts/0    00:00:00 bash  
29545 pts/0    00:00:00 ps  
[root@ip-172-31-45-240 ~]# ps a  
  PID TTY      STAT   TIME COMMAND  
3204 tty1    Ss+    0:00 /sbin/agetty --noclear tty1 linux  
3205 ttys0    Ss+    0:00 /sbin/agetty --keep-baud 115200,38400,9600 ttys0 vt22  
29502 pts/0    Ss    0:00 -bash  
29525 pts/0    S     0:00 sudo -i  
29526 pts/0    S     0:00 -bash  
29546 pts/0    R+    0:00 ps a  
[root@ip-172-31-45-240 ~]#
```



Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...

```
[root@ip-172-31-45-240 ~]# ps ax
```

@thomasdcameron || thomas.cameron@camerontech.com



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...  
2893 ? Ss 0:00 /sbin/dhclient -q -lf /var/lib/dhclient/dhclient--eth  
2994 ? Ss 0:00 /sbin/dhclient -6 -nw -lf /var/lib/dhclient/dhclient6  
3128 ? Ss 0:00 /usr/libexec/postfix/master -w  
3129 ? S 0:00 pickup -l -t unix -u  
3130 ? S 0:00 qmgr -l -t unix -u  
3167 ? Ssl 0:00 /usr/sbin/rsyslogd -n  
3170 ? Ssl 0:00 /usr/bin/amazon-ssm-agent  
3173 ? Ss 0:00 /usr/sbin/sshd -D  
3181 ? Ss 0:00 /usr/sbin/atd -f  
3183 ? Ss 0:00 /usr/sbin/crond -n  
3204 tty1 Ss+ 0:00 /sbin/agetty --noclear tty1 linux  
3205 ttys0 Ss+ 0:00 /sbin/agetty --keep-baud 115200,38400,9600 ttys0 vt22  
3228 ? Ss 0:00 /usr/sbin/acpid  
3324 ? Ss 0:00 /usr/sbin/anacron -s  
3401 ? I 0:00 [kworker/0:1]  
3402 ? I 0:00 [kworker/0:0]  
29484 ? I 0:00 [kworker/0:2]  
29487 ? Ss 0:00 sshd: ec2-user [priv]  
29501 ? S 0:00 sshd: ec2-user@pts/0  
29502 pts/0 Ss 0:00 -bash  
29525 pts/0 S 0:00 sudo -i  
29526 pts/0 S 0:00 -bash  
29547 pts/0 R+ 0:00 ps ax  
[root@ip-172-31-45-240 ~]#
```

A screenshot of a terminal window titled "Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...". The window shows the command "ps axf" being run, which typically lists all processes in a flat tree format. The terminal has a dark purple background and a light gray border. The title bar includes standard window controls (red, yellow, green) and a blue icon.

@thomasdcameron || thomas.cameron@camerontech.com



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...  
2678 ? Ss 0:00 /sbin/rngd -f  
2690 ? S 0:00 /usr/sbin/chronyrd  
2702 ? Ssl 0:00 /usr/sbin/gssproxy -D  
2893 ? Ss 0:00 /sbin/dhclient -q -lf /var/lib/dhclient/dhclient--eth  
2994 ? Ss 0:00 /sbin/dhclient -6 -nw -lf /var/lib/dhclient/dhclient6  
3128 ? Ss 0:00 /usr/libexec/postfix/master -w  
3129 ? S 0:00 \_ pickup -l -t unix -u  
3130 ? S 0:00 \_ qmgr -l -t unix -u  
3167 ? Ssl 0:00 /usr/sbin/rsyslogd -n  
3170 ? Ssl 0:00 /usr/bin/amazon-ssm-agent  
3173 ? Ss 0:00 /usr/sbin/sshd -D  
29487 ? Ss 0:00 \_ sshd: ec2-user [priv]  
29501 ? S 0:00 \_ sshd: ec2-user@pts/0  
29502 pts/0 Ss 0:00 \_ -bash  
29525 pts/0 S 0:00 \_ sudo -i  
29526 pts/0 S 0:00 \_ -bash  
29586 pts/0 R+ 0:00 \_ ps axf  
3181 ? Ss 0:00 /usr/sbin/atd -f  
3183 ? Ss 0:00 /usr/sbin/crond -n  
3204 tty1 Ss+ 0:00 /sbin/agetty --noclear tty1 linux  
3205 ttys0 Ss+ 0:00 /sbin/agetty --keep-baud 115200,38400,9600 ttys0 vt22  
3228 ? Ss 0:00 /usr/sbin/acpid  
3324 ? Ss 0:00 /usr/sbin/anacron -s  
[root@ip-172-31-45-240 ~]#
```

```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@ec2-18-219-197-165.us-east-2.compute.amazonaws.com — 202x56

[[root@ip-172-31-45-240 ~]# ps alx
F  UID  PID  PPID PRI  NI   VSZ  RSS WCHAN STAT TTY      TIME COMMAND
4  0     1    0  20  0  43524  5220 SyS_ep Ss  ?  0:01 /usr/lib/systemd/systemd --switched-root --system --deserialize 22
1  0     2    0  20  0    0  0 - S  ?  0:00 [kthreadd]
1  0     4    2  0 -20  0    0  0 - I<  ?  0:00 [kworker/0:0H]
1  0     5    2  20  0    0  0 - I  ?  0:00 [kworker/u30:0]
1  0     6    2  0 -20  0    0  0 - I<  ?  0:00 [mm_percpu_wq]
1  0     7    2  20  0    0  0 - S  ?  0:00 [ksoftirqd/0]
1  0     8    2  20  0    0  0 rcu_gp I  ?  0:00 [rcu_sched]
1  0     9    2  20  0    0  0 - I  ?  0:00 [rcu_bh]
1  0    10    2 -100 -    0  0 - S  ?  0:00 [migration/0]
5  0    11    2 -100 -    0  0 - S  ?  0:00 [watchdog/0]
1  0    12    2  20  0    0  0 - S  ?  0:00 [cpuhp/0]
5  0    13    2  20  0    0  0 - S  ?  0:00 [kdevtmpfs]
1  0    14    2  0 -20  0    0  0 - I<  ?  0:00 [netns]
1  0    15    2  20  0    0  0 - I  ?  0:00 [kworker/u30:1]
1  0    20    2  20  0    0  0 - S  ?  0:00 [xenbus]
1  0    21    2  20  0    0  0 - S  ?  0:00 [xenwatch]
1  0   171    2  20  0    0  0 watchd S  ?  0:00 [khungtaskd]
1  0   172    2  20  0    0  0 - S  ?  0:00 [oom_reaper]
1  0   173    2  0 -20  0    0  0 - I<  ?  0:00 [writeback]
1  0   175    2  20  0    0  0 - S  ?  0:00 [kcompactd0]
1  0   176    2  25  5    0  0 - SN  ?  0:00 [ksmd]
1  0   177    2  39  19   0  0 - SN  ?  0:00 [khugepaged]
1  0   178    2  0 -20  0    0  0 - I<  ?  0:00 [crypto]
1  0   179    2  0 -20  0    0  0 - I<  ?  0:00 [kintegrityd]
1  0   181    2  0 -20  0    0  0 - I<  ?  0:00 [kblockd]
1  0   534    2  0 -20  0    0  0 - I<  ?  0:00 [md]
1  0   537    2  0 -20  0    0  0 - I<  ?  0:00 [edac-poller]
1  0   542    2  0 -20  0    0  0 - I<  ?  0:00 [watchdogd]
1  0   683    2  20  0    0  0 - S  ?  0:00 [kauditfd]
1  0   689    2  20  0    0  0 - S  ?  0:00 [kswapd0]
1  0   821    2  0 -20  0    0  0 - I<  ?  0:00 [kthrotld]
1  0   871    2  0 -20  0    0  0 - I<  ?  0:00 [kstrp]
1  0   899    2  0 -20  0    0  0 - I<  ?  0:00 [ipv6_addrconf]
1  0  1736    2  0 -20  0    0  0 - I<  ?  0:00 [ata_sff]
1  0  1751    2  20  0    0  0 - S  ?  0:00 [scsi_eh_0]
1  0  1752    2  0 -20  0    0  0 - I<  ?  0:00 [scsi_tmf_0]
1  0  1756    2  20  0    0  0 - S  ?  0:00 [scsi_eh_1]
1  0  1758    2  0 -20  0    0  0 - I<  ?  0:00 [scsi_tmf_1]
1  0  1854    2  0 -20  0    0  0 - I<  ?  0:00 [xfsalloc]
1  0  1855    2  0 -20  0    0  0 - I<  ?  0:00 [xfs_mru_cache]
1  0  1857    2  0 -20  0    0  0 - I<  ?  0:00 [xfs-buf/xvda1]
1  0  1858    2  0 -20  0    0  0 - I<  ?  0:00 [xfs-data/xvda1]
1  0  1859    2  0 -20  0    0  0 - I<  ?  0:00 [xfs-conv/xvda1]
1  0  1860    2  0 -20  0    0  0 - I<  ?  0:00 [xfs-cil/xvda1]
1  0  1861    2  0 -20  0    0  0 - I<  ?  0:00 [xfs-reclaim/xvd]
1  0  1862    2  0 -20  0    0  0 - I<  ?  0:00 [xfs-log/xvda1]
1  0  1863    2  0 -20  0    0  0 - I<  ?  0:00 [xfs-eofblocks/x]
1  0  1864    2  20  0    0  0 kthrea S  ?  0:00 [xfsaild/xvda1]
1  0  1865    2  0 -20  0    0  0 - I<  ?  0:00 [kworker/0:1H]
4  0  1931    1  20  0  49308 10964 SyS_ep Ss  ?  0:00 /usr/lib/systemd/systemd-journald
1  0  1943    2  0 -20  0    0  0 - I<  ?  0:00 [ena]
4  0  1956    1  20  0  116824 2132 core_s Ss  ?  0:00 /usr/sbin/lvmtd -f
4  0  2084    1  20  0  46036 4080 SyS_ep Ss  ?  0:00 /usr/lib/systemd/systemd-udevd
1  0  2637    2  0 -20  0    0  0 - I<  ?  0:00 [rpcliod]
```

Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@ec2-18-219-197-165.us-east-2.compute.amazonaws.com — 202x56												
F	UID	PID	PPID	PRI	NI	VSZ	RSS	WCHAN	STAT	TTY	TIME	COMMAND
1	0	2	0	20	0	0	0	-	S	?	0:00	[kthreadd]
1	0	4	2	0	-20	0	0	-	I<	?	0:00	\_ [kworker/0:0H]
1	0	5	2	20	0	0	0	-	I	?	0:00	\_ [kworker/u30:0]
1	0	6	2	0	-20	0	0	-	I<	?	0:00	\_ [mm_percpu_wq]
1	0	7	2	20	0	0	0	-	S	?	0:00	\_ [ksoftirqd/0]
1	0	8	2	20	0	0	0	rcu_gp	I	?	0:00	\_ [rcu_sched]
1	0	9	2	20	0	0	0	-	I	?	0:00	\_ [rcu_bh]
1	0	10	2	-100	-	0	0	-	S	?	0:00	\_ [migration/0]
5	0	11	2	-100	-	0	0	-	S	?	0:00	\_ [watchdog/0]
1	0	12	2	20	0	0	0	-	S	?	0:00	\_ [cpuhp/0]
5	0	13	2	20	0	0	0	-	S	?	0:00	\_ [kdevtmpfs]
1	0	14	2	0	-20	0	0	-	I<	?	0:00	\_ [netns]
1	0	15	2	20	0	0	0	-	I	?	0:00	\_ [kworker/u30:1]
1	0	20	2	20	0	0	0	-	S	?	0:00	\_ [xenbus]
1	0	21	2	20	0	0	0	-	S	?	0:00	\_ [xenwatch]
1	0	171	2	20	0	0	0	watchd	S	?	0:00	\_ [khungtaskd]
1	0	172	2	20	0	0	0	-	S	?	0:00	\_ [oom_reaper]
1	0	173	2	0	-20	0	0	-	I<	?	0:00	\_ [writeback]
1	0	175	2	20	0	0	0	-	S	?	0:00	\_ [kcompactd0]
1	0	176	2	25	5	0	0	-	SN	?	0:00	\_ [ksmd]
1	0	177	2	39	19	0	0	-	SN	?	0:00	\_ [khugepaged]
1	0	178	2	0	-20	0	0	-	I<	?	0:00	\_ [crypto]
1	0	179	2	0	-20	0	0	-	I<	?	0:00	\_ [kintegrityd]
1	0	181	2	0	-20	0	0	-	I<	?	0:00	\_ [kblockd]
1	0	534	2	0	-20	0	0	-	I<	?	0:00	\_ [md]
1	0	537	2	0	-20	0	0	-	I<	?	0:00	\_ [edac-poller]
1	0	542	2	0	-20	0	0	-	I<	?	0:00	\_ [watchdogd]
1	0	683	2	20	0	0	0	-	S	?	0:00	\_ [kauditfd]
1	0	689	2	20	0	0	0	-	S	?	0:00	\_ [kswapd0]
1	0	821	2	0	-20	0	0	-	I<	?	0:00	\_ [kthrotld]
1	0	871	2	0	-20	0	0	-	I<	?	0:00	\_ [kstrp]
1	0	899	2	0	-20	0	0	-	I<	?	0:00	\_ [ipv6_addrconf]
1	0	1736	2	0	-20	0	0	-	I<	?	0:00	\_ [ata_sff]
1	0	1751	2	20	0	0	0	-	S	?	0:00	\_ [scsi_eh_0]
1	0	1752	2	0	-20	0	0	-	I<	?	0:00	\_ [scsi_tmf_0]
1	0	1756	2	20	0	0	0	-	S	?	0:00	\_ [scsi_eh_1]
1	0	1758	2	0	-20	0	0	-	I<	?	0:00	\_ [scsi_tmf_1]
1	0	1854	2	0	-20	0	0	-	I<	?	0:00	\_ [xfsalloc]
1	0	1855	2	0	-20	0	0	-	I<	?	0:00	\_ [xfs_mru_cache]
1	0	1857	2	0	-20	0	0	-	I<	?	0:00	\_ [xfs-buf/xvda1]
1	0	1858	2	0	-20	0	0	-	I<	?	0:00	\_ [xfs-data/xvda1]
1	0	1859	2	0	-20	0	0	-	I<	?	0:00	\_ [xfs-conv/xvda1]
1	0	1860	2	0	-20	0	0	-	I<	?	0:00	\_ [xfs-cil/xvda1]
1	0	1861	2	0	-20	0	0	-	I<	?	0:00	\_ [xfs-reclaim/xvd]
1	0	1862	2	0	-20	0	0	-	I<	?	0:00	\_ [xfs-log/xvda1]
1	0	1863	2	0	-20	0	0	-	I<	?	0:00	\_ [xfs-eofblocks/x]
1	0	1864	2	20	0	0	0	kthrea	S	?	0:00	\_ [xfsaidd/xvda1]
1	0	1865	2	0	-20	0	0	-	I<	?	0:00	\_ [kworker/0:1H]
1	0	1943	2	0	-20	0	0	-	I<	?	0:00	\_ [enal]
1	0	2637	2	0	-20	0	0	-	I<	?	0:00	\_ [rpciod]
1	0	2639	2	0	-20	0	0	-	I<	?	0:00	\_ [xpriod]
1	0	3402	2	20	0	0	0	-	I	?	0:00	\_ [kworker/0:0]
1	0	29484	2	20	0	0	0	-	I	?	0:00	\_ [kworker/0:2]
1	0	29587	2	20	0	0	0	-	I	?	0:00	\_ [kworker/0:1]

Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@ec2-18-219-197-165.us-east-2.compute.amazonaws.com — 202x56

```
1 0 899 2 0 -20 0 0 - I< ? 0:00 \_ [ipv6_addrconf]
1 0 1736 2 0 -20 0 0 - I< ? 0:00 \_ [ata_sff]
1 0 1751 2 20 0 0 0 - S ? 0:00 \_ [scsi_eh_0]
1 0 1752 2 0 -20 0 0 - I< ? 0:00 \_ [scsi_tmf_0]
1 0 1756 2 20 0 0 0 - S ? 0:00 \_ [scsi_eh_1]
1 0 1758 2 0 -20 0 0 - I< ? 0:00 \_ [scsi_tmf_1]
1 0 1854 2 0 -20 0 0 - I< ? 0:00 \_ [xfsalloc]
1 0 1855 2 0 -20 0 0 - I< ? 0:00 \_ [xfs_mru_cache]
1 0 1857 2 0 -20 0 0 - I< ? 0:00 \_ [xfs-buf/xvda1]
1 0 1858 2 0 -20 0 0 - I< ? 0:00 \_ [xfs-data/xvda1]
1 0 1859 2 0 -20 0 0 - I< ? 0:00 \_ [xfs-conv/xvda1]
1 0 1860 2 0 -20 0 0 - I< ? 0:00 \_ [xfs-cil/xvda1]
1 0 1861 2 0 -20 0 0 - I< ? 0:00 \_ [xfs-reclaim/xvd]
1 0 1862 2 0 -20 0 0 - I< ? 0:00 \_ [xfs-log/xvda1]
1 0 1863 2 0 -20 0 0 - I< ? 0:00 \_ [xfs-eofblocks/x]
1 0 1864 2 20 0 0 0 kthrea S ? 0:00 \_ [xfsaild/xvda1]
1 0 1865 2 0 -20 0 0 - I< ? 0:00 \_ [kworker/0:1H]
1 0 1943 2 0 -20 0 0 - I< ? 0:00 \_ [ena]
1 0 2637 2 0 -20 0 0 - I< ? 0:00 \_ [rpciod]
1 0 2639 2 0 -20 0 0 - I< ? 0:00 \_ [xpriod]
1 0 3402 2 20 0 0 0 - I ? 0:00 \_ [kworker/0:0]
1 0 29484 2 20 0 0 0 - I ? 0:00 \_ [kworker/0:2]
1 0 29587 2 20 0 0 0 - I ? 0:00 \_ [kworker/0:1]
4 0 1 0 20 0 43524 5220 SyS_ep Ss ? 0:01 /usr/lib/systemd/systemd --switched-root --system --deserialize 22
4 0 1931 1 20 0 49308 10964 SyS_ep Ss ? 0:00 /usr/lib/systemd/systemd-journald
4 0 1956 1 20 0 116824 2132 core_s Ss ? 0:00 /usr/sbin/lvmetad -f
4 0 2084 1 20 0 46036 4080 SyS_ep Ss ? 0:00 /usr/lib/systemd/systemd-udevd
5 0 2642 1 16 -4 64328 2328 SyS_ep S<sl ? 0:00 /sbin/auditd
4 81 2666 1 20 0 60420 4012 SyS_ep Ss ? 0:00 /usr/bin/dbus-daemon --system --address=systemd: --nofork --nopidfile --systemd-activation
5 32 2667 1 20 0 73828 3564 SyS_po Ss ? 0:00 /sbin/rpcbind -w
4 0 2668 1 20 0 28512 2968 SyS_ep Ss ? 0:00 /usr/lib/systemd/systemd-logind
4 999 2671 1 20 0 12656 1756 core_s Ss ? 0:00 /sbin/lsmd -d
4 0 2678 1 20 0 13144 880 SyS_po Ss ? 0:00 /sbin/rngd -f
5 997 2690 1 20 0 122644 4072 core_s S ? 0:00 /usr/sbin/chronyd
1 0 2702 1 20 0 208356 3628 SyS_ep Ssl ? 0:00 /usr/sbin/gssproxy -D
5 0 2893 1 20 0 105448 4096 core_s Ss ? 0:00 /sbin/dhclient -q -lf /var/lib/dhclient/dhclient--eth0.lease -pf /var/run/dhclient-eth0.pid -H ip-172-31-45-240 eth0
1 0 2994 1 20 0 105448 3856 core_s Ss ? 0:00 /sbin/dhclient -6 -nw -lf /var/lib/dhclient/dhclient6--eth0.lease -pf /var/run/dhclient6-eth0.pid eth0 -H ip-172-31-45-240
5 0 3128 1 20 0 90416 4740 SyS_ep Ss ? 0:00 /usr/libexec/postfix/master -w
4 89 3129 3128 20 0 90504 6828 SyS_ep S ? 0:00 \_ pickup -l -t unix -u
4 89 3130 3128 20 0 90580 6652 SyS_ep S ? 0:00 \_ qmgr -l -t unix -u
4 0 3167 1 20 0 224824 5896 core_s Ssl ? 0:00 /usr/sbin/rsyslogd -n
4 0 3170 1 20 0 542424 12992 SyS_ep Ssl ? 0:00 /usr/bin/amazon-ssm-agent
4 0 3173 1 20 0 112924 7192 core_s Ss ? 0:00 /usr/sbin/sshd -D
4 0 29487 3173 20 0 150608 8876 SyS_po Ss ? 0:00 \_ sshd: ec2-user [priv]
5 1000 29501 29487 20 0 150608 4660 core_s S ? 0:00 \_ sshd: ec2-user@pts/0
0 1000 29502 29501 20 0 124792 3780 - Ss pts/0 0:00 \_ -bash
4 0 29525 29502 20 0 216936 6344 SyS_po S pts/0 0:00 \_ sudo -i
4 0 29526 29525 20 0 124792 4064 - S pts/0 0:00 \_ -bash
0 0 29594 29526 20 0 160376 3896 - R+ pts/0 0:00 \_ ps alxf
4 0 3181 1 20 0 27916 2236 hrtime Ss ? 0:00 /usr/sbin/atd -f
4 0 3183 1 20 0 135116 3128 hrtime Ss ? 0:00 /usr/sbin/crond -n
4 0 3204 1 20 0 121336 1760 core_s Ss+ tty1 0:00 /sbin/agetty --noclear tty1 linux
4 0 3205 1 20 0 120984 2128 core_s Ss+ ttyS0 0:00 /sbin/agetty --keep-baud 115200,38400,9600 ttyS0 vt220
1 0 3228 1 20 0 4308 108 core_s Ss ? 0:00 /usr/sbin/acpid
1 0 3324 1 20 0 134212 2724 - Ss ? 0:00 /usr/sbin/anacron -s
[root@ip-172-31-45-240 ~]#
```

# Tools for baseline

- **Check out what's running, how much memory, cpu, etc.**

# Tools for baseline

- **Check out what's running, how much memory, cpu, etc.**
- **Record this info so you can compare later**

# Tools for baseline

- **free**



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...

[[root@ip-172-31-45-240 ~]# free
      total        used        free      shared  buff/cache   available
Mem:       1007348      72132      748724          428      186492      783388
Swap:            0          0          0
[root@ip-172-31-45-240 ~]# ]
```

# Tools for baseline

- **top**



Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...

```
[root@ip-172-31-45-240 ~]# top
```

@thomasdcameron || thomas.cameron@camerontech.com

# TXLF 2019



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...
```

```
top - 12:33:46 up 1:03, 1 user, load average: 0.00, 0.00, 0.00
Tasks: 86 total, 1 running, 49 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.0 us, 0.0 sy, 0.0 ni, 100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 1007348 total, 724448 free, 61588 used, 221312 buff/cache
KiB Swap: 0 total, 0 free, 0 used. 787852 avail Mem

PID USER PR NI VIRT RES SHR S %CPU %MEM TIME+ COMMAND
 1 root 20 0 43524 5220 3916 S 0.0 0.5 0:01.52 systemd
 2 root 20 0 0 0 0 S 0.0 0.0 0:00.00 kthreadd
 4 root 0 -20 0 0 0 I 0.0 0.0 0:00.00 kworker/0:0H
 5 root 20 0 0 0 0 I 0.0 0.0 0:00.01 kworker/u30+
 6 root 0 -20 0 0 0 I 0.0 0.0 0:00.00 mm_percpu_wq
 7 root 20 0 0 0 0 S 0.0 0.0 0:00.10 ksoftirqd/0
 8 root 20 0 0 0 0 I 0.0 0.0 0:00.13 rcu_sched
 9 root 20 0 0 0 0 I 0.0 0.0 0:00.00 rcu_bh
10 root rt 0 0 0 0 S 0.0 0.0 0:00.00 migration/0
11 root rt 0 0 0 0 S 0.0 0.0 0:00.00 watchdog/0
12 root 20 0 0 0 0 S 0.0 0.0 0:00.00 cpuhp/0
13 root 20 0 0 0 0 S 0.0 0.0 0:00.00 kdevtmpfs
14 root 0 -20 0 0 0 I 0.0 0.0 0:00.00 netns
15 root 20 0 0 0 0 I 0.0 0.0 0:00.02 kworker/u30+
20 root 20 0 0 0 0 S 0.0 0.0 0:00.00 xenbus
21 root 20 0 0 0 0 S 0.0 0.0 0:00.01 xenwatch
171 root 20 0 0 0 0 S 0.0 0.0 0:00.00 khungtaskd
```

# Tools for baseline

- **top**
  - **t** to change summary info
  - **m** to change memory summary
  - **r** to renice
  - **k** to kill

# Tools for baseline

- **top**
  - **Sometimes your system will seem fine but the app is really slow**

# Tools for baseline

- **top**
  - **Sometimes your system will seem fine but the app is really slow**
  - **Looking at a per-cpu view can explain why**

# Tools for baseline

- **top**
  - **Sometimes your system will seem fine but the app is really slow**
  - **Looking at a per-cpu view can explain why**
  - **Press 1**



```
tdcam — root@wintermute:~ — ssh root@wintermute — 80x24
top - 07:37:13 up 7:36, 1 user, load average: 0.05, 0.04, 0.05
Tasks: 347 total, 1 running, 346 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.1 us, 0.1 sy, 0.0 ni, 99.8 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 65962352 total, 58322184 free, 5099412 used, 2540756 buff/cache
KiB Swap: 2150396 total, 2150396 free, 0 used. 60309616 avail Mem

PID USER PR NI VIRT RES SHR S %CPU %MEM TIME+ COMMAND
8486 qemu 20 0 5320900 4.2g 11164 S 4.7 6.7 28:52.22 qemu-kvm
6676 root 20 0 162144 2500 1592 R 0.7 0.0 0:00.07 top
 10 root 20 0 0 0 0 S 0.3 0.0 0:06.17 rcu_sched
 7402 root 20 0 228156 11156 6516 S 0.3 0.0 0:16.86 snmpd
 8025 root 20 0 50608 2048 672 S 0.3 0.0 0:11.40 cmaperfd
11572 root 20 0 0 0 0 S 0.3 0.0 0:03.50 iscsi_ttx
11575 root 20 0 0 0 0 S 0.3 0.0 0:02.53 iscsi_trx
11576 root 20 0 0 0 0 S 0.3 0.0 0:25.46 iscsi_ttx
 1 root 20 0 195400 8556 4184 S 0.0 0.0 0:07.77 systemd
 2 root 20 0 0 0 0 S 0.0 0.0 0:00.03 kthreadd
 3 root 20 0 0 0 0 S 0.0 0.0 0:00.12 ksoftirqd/0
 5 root 0 -20 0 0 0 S 0.0 0.0 0:00.00 kworker/0:+
 8 root rt 0 0 0 0 S 0.0 0.0 0:00.15 migration/0
 9 root 20 0 0 0 0 S 0.0 0.0 0:00.00 rcu_bh
11 root 0 -20 0 0 0 S 0.0 0.0 0:00.00 lru-add-dr+
12 root rt 0 0 0 0 S 0.0 0.0 0:00.08 watchdog/0
13 root rt 0 0 0 0 S 0.0 0.0 0:00.07 watchdog/1
```

tdcam — root@wintermute:~ — ssh root@wintermute — 97x39

```
top - 07:38:50 up 7:38, 1 user, load average: 0.11, 0.06, 0.05
Tasks: 348 total, 1 running, 347 sleeping, 0 stopped, 0 zombie
%Cpu0 : 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu1 : 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu2 : 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu3 : 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu4 : 0.0 us, 0.3 sy, 0.0 ni, 99.7 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu5 : 0.0 us, 0.3 sy, 0.0 ni, 99.7 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu6 : 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu7 : 0.0 us, 0.3 sy, 0.0 ni, 99.7 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu8 : 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu9 : 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu10: 0.3 us, 0.0 sy, 0.0 ni, 99.7 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu11: 1.0 us, 0.0 sy, 0.0 ni, 99.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu12: 0.0 us, 0.3 sy, 0.0 ni, 99.7 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu13: 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu14: 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu15: 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu16: 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu17: 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu18: 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu19: 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu20: 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu21: 1.0 us, 0.0 sy, 0.0 ni, 99.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu22: 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu23: 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 65962352 total, 58320960 free, 5099960 used, 2541432 buff/cache
KiB Swap: 2150396 total, 2150396 free, 0 used. 60308464 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
8486	qemu	20	0	5304508	4.2g	11164	S	4.0	6.7	28:57.00	qemu-kvm
7933	root	20	0	55392	3164	1952	S	0.7	0.0	0:19.72	cmahostd
<b>7280</b>	<b>root</b>	<b>20</b>	<b>0</b>	<b>162152</b>	<b>2520</b>	<b>1592</b>	<b>R</b>	<b>0.3</b>	<b>0.0</b>	<b>0:00.07</b>	<b>top</b>
11571	root	20	0	0	0	0	S	0.3	0.0	0:01.34	iscsi_trx
1	root	20	0	195400	8556	4184	S	0.0	0.0	0:07.79	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.03	kthreadd
3	root	20	0	0	0	0	S	0.0	0.0	0:00.12	ksoftirqd/0
5	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	kworker/0:0H
8	root	rt	0	0	0	0	S	0.0	0.0	0:00.15	migration/0

# Tools for baseline

- **vmstat**
  - Part of the **procps-ng** package

# Tools for baseline

- **vmstat**
  - Part of the procps-ng package
  - **vmstat [number]**

# Tools for baseline

- **vmstat**
  - Part of the procps-ng package
  - **vmstat [number]**
  - **Great for seeing what's going on with memory/cpu**



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...
[[root@ip-172-31-45-240 ~]# vmstat
procs -----memory----- --swap-- -----io---- -system-- -----cpu-----
r b    swpd   free   buff   cache   si   so    bi    bo   in   cs us sy id wa st
0 0      0 724696  2088 219284    0    0    44     7   20   63  0  0 99  0  0
[root@ip-172-31-45-240 ~]# ]
```



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@ec2-18-21...  
[root@ip-172-31-45-240 ~]# vmstat 3  
procs -----memory----- --swap-- -----io---- -system-- -----cpu-----  
r b swpd free buff cache si so bi bo in cs us sy id wa st  
0 0 0 724704 2088 219284 0 0 43 7 20 62 0 0 99 0 0  
0 0 0 724696 2088 219284 0 0 0 0 10 15 0 0 100 0 0  
0 0 0 724696 2088 219284 0 0 0 0 5 8 0 0 100 0 0  
0 0 0 724696 2088 219284 0 0 0 0 6 8 0 0 100 0 0  
0 0 0 724696 2088 219284 0 0 0 1 7 12 0 0 100 0 0  
0 0 0 724696 2088 219284 0 0 0 0 5 8 0 0 100 0 0  
0 0 0 724696 2088 219284 0 0 0 44 13 11 0 0 100 0 0  
0 0 0 724696 2088 219284 0 0 0 0 7 11 0 0 100 0 0  
0 0 0 724696 2088 219284 0 0 0 0 5 8 0 0 100 0 0  
0 0 0 722944 2088 219328 0 0 0 0 13 27 0 0 100 0 0  
1 0 0 724200 2088 219360 0 0 5 0 20 47 0 0 100 0 0  
0 0 0 724200 2088 219360 0 0 0 0 6 10 0 0 100 0 0  
0 0 0 724200 2088 219360 0 0 0 0 8 13 0 0 100 0 0  
0 0 0 724136 2088 219424 0 0 0 55 20 35 0 0 100 0 0
```



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@ec2-18-21...  
[root@ip-172-31-45-240 ~]# vmstat -d  
disk- -----reads----- -----writes----- -----IO-----  
          total merged sectors      ms      total merged sectors      ms      cur      sec  
xvda    12902        14   393363    7404     3961        170   64170    5988        0         6  
[root@ip-172-31-45-240 ~]#
```

# TXLF 2019



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@ec2-18-21...

[[root@ip-172-31-45-240 ~]# vmstat -m
Cache           Num  Total   Size  Pages
rpc_inode_cache      23     23    704    23
xfs_dqtrx          0      0    528    31
xfs_rui_item        0      0    672    24
xfs_rud_item        0      0    152    26
xfs_ili            264    264   168    24
xfs_inode          15096  15096   960    34
xfs_efd_item        78     78   416    39
xfs_buf_item         66     66   248    33
xfs_da_state        34     34   480    34
xfs_btree_cur       35     35   232    35
xfs_log_ticket      44     44   184    22
bio-2              200    200   320    25
kcopyd_job          0      0  3312     9
dm_uevent          0      0  2632    12
dax_cache           21     21   768    21
RAWv6              28     28  1152    28
UDPV6              25     25  1280    25
tw_sock_TCPv6       0      0   248    33
request_sock_TCPv6 0      0   312    26
TCPv6              15     15  2176    15
cfq_io_cq           0      0   120    34
Cache           Num  Total   Size  Pages
```

# Tools for baseline

- **iostat**

# Tools for baseline

- **iostat**
  - Part of the sysstat package



```
● ○ ● Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@ec2-18-21...
[[root@ip-172-31-45-240 ~]# iostat
Linux 4.14.121-109.96.amzn2.x86_64 (ip-172-31-45-240.us-east-2.compute.internal) 0
6/01/2019      _x86_64_      (1 CPU)

avg-cpu: %user  %nice %system %iowait  %steal  %idle
          0.12    0.08    0.18    0.08    0.19   99.35

Device:         tps    kB_read/s    kB_wrtn/s    kB_read    kB_wrtn
xvda           3.22     37.27       6.22    196933     32841

[root@ip-172-31-45-240 ~]#
```

A screenshot of a terminal window titled "Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...". The window displays the output of the "iostat -p" command on a Linux system. The output shows CPU usage statistics and disk I/O metrics for devices xvda and xvdap1.

```
[root@ip-172-31-45-240 ~]# iostat -p
Linux 4.14.121-109.96.amzn2.x86_64 (ip-172-31-45-240.us-east-2.compute.internal)
06/01/2019      _x86_64_      (1 CPU)

avg-cpu:  %user   %nice  %system  %iowait  %steal   %idle
          0.12    0.08    0.17    0.07    0.19    99.37

Device:            tps    kB_read/s    kB_wrtn/s    kB_read    kB_wrtn
xvda              3.15     35.84         6.26    196933     34386
xvdap1             3.12     35.63         6.26    195781     34386

[root@ip-172-31-45-240 ~]#
```

# Tools for baseline

- **sar**

# Tools for baseline

- **sar**
  - **Part of the sysstat package**

# Tools for baseline

- **sar**
  - **Part of the sysstat package**
  - **Finer grained view over time**



```
[root@ip-172-31-45-240:~]— ssh -i ThomassKeyPair.pem ec2-user@ec2-18-219...
[[root@ip-172-31-45-240 ~]# sar
Linux 4.14.114-105.126.amzn2.x86_64 (ip-172-31-45-240.us-east-2.compute.internal
)       06/01/2019      _x86_64_          (1 CPU)

05:11:49 AM      LINUX RESTART

05:18:06 AM      LINUX RESTART

05:20:01 AM      CPU      %user      %nice      %system      %iowait      %steal      %idle
05:30:01 AM      all       0.04       0.00       0.03        0.01        0.01      99.91
Average:         all       0.04       0.00       0.03        0.01        0.01      99.91

11:30:48 AM      LINUX RESTART

11:40:01 AM      CPU      %user      %nice      %system      %iowait      %steal      %idle
11:50:01 AM      all       0.00       0.00       0.00        0.00        0.00      99.99
12:00:01 PM      all       0.01       0.00       0.01        0.00        0.00      99.98
12:10:01 PM      all       0.10       0.00       0.03        0.01        0.01      99.85
12:20:01 PM      all       0.00       0.74       1.05        0.48        0.03      97.69
12:30:01 PM      all       0.02       0.00       0.02        0.00        0.01      99.94
12:40:01 PM      all       0.01       0.00       0.01        0.00        0.00      99.97
12:50:01 PM      all       0.01       0.00       0.01        0.00        0.00      99.97
01:00:01 PM      all       0.04       0.00       0.03        0.01        0.01      99.92
Average:         all       0.03       0.09       0.14        0.06        0.01      99.67
```

# Tools for baseline

- **ss**

# Tools for baseline

- **ss**
  - Network information

# Tools for baseline

- **ss**
  - **Network information**
  - **Replaces older utils like netstat**

# TXLF 2019



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...  
[[root@ip-172-31-45-240 ~]# ss  
NetidState Recv-Q Send-Q Local Address:Port Peer Address:Port  
u_strESTAB 0 0 * 18653 * 18652  
u_strESTAB 0 0 * 19312 * 19313  
u_strESTAB 0 0 /run/systemd/journal/stdout 16262 * 16261  
u_strESTAB 0 0 * 16061 * 16062  
u_strESTAB 0 0 * 18643 * 18644  
u_strESTAB 0 0 * 18646 * 18647  
u_strESTAB 0 0 /run/systemd/journal/stdout 16807 * 16806  
u_strESTAB 0 0 * 16331 * 16332  
u_strESTAB 0 0 * 16261 * 16262  
u_strESTAB 0 0 * 16806 * 16807  
u_strESTAB 0 0 * 18644 * 18643  
u_strESTAB 0 0 /run/dbus/system_bus_socket 16333 * 16211  
u_strESTAB 0 0 /run/systemd/journal/stdout 14900 * 14899  
u_strESTAB 0 0 * 14899 * 14900  
u_strESTAB 0 0 * 16647 * 16648  
u_strESTAB 0 0 /run/dbus/system_bus_socket 16654 * 16653  
u_strESTAB 0 0 * 16332 * 16331  
u_strESTAB 0 0 * 15027 * 15174  
u_strESTAB 0 0 * 50412 * 50413  
u_strESTAB 0 0 * 16211 * 16333  
u_strESTAB 0 0 * 16062 * 16061  
u_strESTAB 0 0 /run/systemd/journal/stdout 15174 * 15027
```

# TXLF 2019



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...  
[root@ip-172-31-45-240 ~]# ss -t  
State Recv-Q Send-Q Local Address:Port Peer Address:Port  
ESTAB 0 36 172.31.45.240:ssh 68.203.23.181:52617  
ESTAB 0 0 172.31.45.240:ssh 68.203.23.181:52583  
[root@ip-172-31-45-240 ~]# ss -t -a  
State Recv-Q Send-Q Local Address:Port Peer Address:Port  
LISTEN 0 128 0.0.0.0:sunrpc 0.0.0.0:*  
LISTEN 0 128 0.0.0.0:ssh 0.0.0.0:  
LISTEN 0 100 127.0.0.1:smtp 0.0.0.0:  
ESTAB 0 36 172.31.45.240:ssh 68.203.23.181:52617  
ESTAB 0 0 172.31.45.240:ssh 68.203.23.181:52583  
LISTEN 0 128 [:]:sunrpc [:]*  
LISTEN 0 128 [:]:ssh [:]*  
[root@ip-172-31-45-240 ~]#
```

# Tools for baseline

- **iotop**

# Tools for baseline

- **iotop**
  - Watch top i/o loads



Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...

		Total DISK READ :	0.00 B/s		Total DISK WRITE :	0.00 B/s	
		Actual DISK READ:	0.00 B/s		Actual DISK WRITE:	0.00 B/s	
TID	PRIOR	USER	DISK READ	DISK WRITE	SWAPIN	IO>	COMMAND
1	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	systemd ~rialize 22
2	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[kthreadd]
171	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[khungtaskd]
4	be/0	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[kworker/0:0H]
5	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[kworker/u30:0]
6	be/0	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[mm_percpu_wq]
7	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[ksoftirqd/0]
8	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[rcu_sched]
9	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[rcu_bh]
10	rt/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[migration/0]
11	rt/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[watchdog/0]
12	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[cpuhp/0]
13	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[kdevtmpfs]
14	be/0	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[netns]
20	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[xenbus]
21	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[xenwatch]
534	be/0	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[md]
537	be/0	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[edac-poller]
542	be/0	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[watchdogd]
2084	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	systemd-udevd
177	be/7	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	[khugepaged]

# Tools for baseline

- I/O scheduler

# Tools for baseline

- **I/O scheduler**
  - **CFQ (cfq): the default scheduler for many Linux distributions; it places synchronous requests, submitted by processes, into a number of per-process queues and then allocates timeslices for each of the queues to access the disk.**

# Tools for baseline

- **I/O scheduler**
  - **Noop scheduler (noop): the simplest I/O scheduler for the Linux kernel based on the First In First Out (FIFO) queue concept. This scheduler is best suited for SSDs.**

# Tools for baseline

- **I/O scheduler**
  - **Deadline scheduler (deadline): attempts to guarantee a start service time for a request.**

A screenshot of a terminal window titled "Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...". The terminal shows the command "cat /sys/block/xvda/queue/scheduler" being run, which outputs "[noop] cfq".

```
[root@ip-172-31-45-240 ~]# cat /sys/block/xvda/queue/scheduler
[noop] cfq
[root@ip-172-31-45-240 ~]#
```

# Tools for baseline

- CPU management

A screenshot of a terminal window titled 'Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...'. The window shows a single line of text: '[root@ip-172-31-45-240 ~]# yum install cpupowerutils'.

@thomasdcameron || thomas.cameron@camerontech.com

A screenshot of a terminal window titled "Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...". The window shows the following command being run and its output:

```
[root@ip-172-31-45-240 ~]# cpupower frequency-info --governors
analyzing CPU 0:
available cpufreq governors: Not Available
[root@ip-172-31-45-240 ~]#
```



# Tools for tuning

# Tools for tuning

- **I'm lazy**

# Tools for tuning

- **I'm lazy**
- **I like tuned (pronounced “tune-dee”)**

## 3.1. Tuned Overview

**Tuned** is a daemon that uses `udev` to monitor connected devices and statically and dynamically tunes system settings according to a selected profile. **Tuned** is distributed with a number of predefined profiles for common use cases like high throughput, low latency, or powersave. It is possible to modify the rules defined for each profile and customize how to tune a particular device. To revert all changes made to the system settings by a certain profile, you can either switch to another profile or deactivate the **tuned** service.

[https://access.redhat.com/documentation/en-us/red\\_hat\\_enterprise\\_linux/7/html/performance\\_tuning\\_guide/chap-red\\_hat\\_enterprise\\_linux-performance\\_tuning\\_guide-tuned#ch-Tuned-overview](https://access.redhat.com/documentation/en-us/red_hat_enterprise_linux/7/html/performance_tuning_guide/chap-red_hat_enterprise_linux-performance_tuning_guide-tuned#ch-Tuned-overview)



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...  
[root@ip-172-31-45-240 ~]# yum list tuned*  
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd  
Available Packages  
tuned.noarch 2.8.0-5.amzn2.0.1 amzn2-core  
tuned-gtk.noarch 2.8.0-5.amzn2.0.1 amzn2-core  
tuned-profiles-atomic.noarch 2.8.0-5.amzn2.0.1 amzn2-core  
tuned-profiles-compat.noarch 2.8.0-5.amzn2.0.1 amzn2-core  
tuned-profiles-cpu-partitioning.noarch 2.8.0-5.amzn2.0.1 amzn2-core  
tuned-profiles-nfv.noarch 2.8.0-5.amzn2.0.1 amzn2-core  
tuned-profiles-nfv-guest.noarch 2.8.0-5.amzn2.0.1 amzn2-core  
tuned-profiles-nfv-host.noarch 2.8.0-5.amzn2.0.1 amzn2-core  
tuned-profiles-oracle.noarch 2.8.0-5.amzn2.0.1 amzn2-core  
tuned-profiles-realtime.noarch 2.8.0-5.amzn2.0.1 amzn2-core  
tuned-profiles-sap.noarch 2.8.0-5.amzn2.0.1 amzn2-core  
tuned-profiles-sap-hana.noarch 2.8.0-5.amzn2.0.1 amzn2-core  
tuned-utils.noarch 2.8.0-5.amzn2.0.1 amzn2-core  
tuned-utils-systemtap.noarch 2.8.0-5.amzn2.0.1 amzn2-core  
[root@ip-172-31-45-240 ~]#
```



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...

=====
Installing:
tuned           noarch  2.8.0-5.amzn2.0.1      amzn2-core   234 k
Installing for dependencies:
dbus-glib        x86_64  0.100-7.2.amzn2       amzn2-core   103 k
dbus-python       x86_64  1.1.1-9.amzn2.0.2     amzn2-core   206 k
gobject-introspection x86_64  1.50.0-1.amzn2.0.2     amzn2-core   239 k
mozjs17          x86_64  17.0.0-20.amzn2.0.1    amzn2-core   1.4 M
polkit            x86_64  0.112-18.amzn2.1      amzn2-core   169 k
polkit-pkla-compat x86_64  0.1-4.amzn2.0.2      amzn2-core   39 k
python-decorator  noarch  3.4.0-3.amzn2       amzn2-core   27 k
python-gobject-base x86_64  3.22.0-1.amzn2.1     amzn2-core   293 k
python-linux-procfs noarch  0.4.9-3.amzn2       amzn2-core   33 k
python-perf         x86_64  4.14.121-109.96.amzn2  amzn2-core   114 k
python-pyudev       noarch  0.15-9.amzn2       amzn2-core   55 k
python-schedutils  x86_64  0.4-6.amzn2.0.2      amzn2-core   21 k

Transaction Summary
=====
Install 1 Package (+12 Dependent packages)

Total download size: 2.9 M
Installed size: 8.9 M
Is this ok [y/d/N]: 
```

A screenshot of a terminal window titled "Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...". The terminal shows the command: [root@ip-172-31-45-240 ~]# yum install tuned-profiles\* tuned-utils\*  
The terminal window has a light gray background and a dark gray border. The title bar contains the session information and the command being run.

@thomasdcameron || thomas.cameron@camerontech.com



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...

polkit x86_64 0.112-18.amzn2.1 amzn2-core 169 k
polkit-pkla-compat x86_64 0.1-4.amzn2.0.2 amzn2-core 39 k
powertop x86_64 2.3-12.amzn2.0.1 amzn2-core 199 k
python-decorator noarch 3.4.0-3.amzn2 amzn2-core 27 k
python-ethtool x86_64 0.8-5.amzn2.0.2 amzn2-core 33 k
python-gobject-base x86_64 3.22.0-1.amzn2.1 amzn2-core 293 k
python-linux-procfs noarch 0.4.9-3.amzn2 amzn2-core 33 k
python-perf x86_64 4.14.121-109.96.amzn2 amzn2-core 114 k
python-pyudev noarch 0.15-9.amzn2 amzn2-core 55 k
python-schedutils x86_64 0.4-6.amzn2.0.2 amzn2-core 21 k
systemtap x86_64 3.2-4.amzn2.0.2 amzn2-core 142 k
systemtap-client x86_64 3.2-4.amzn2.0.2 amzn2-core 3.7 M
systemtap-devel x86_64 3.2-4.amzn2.0.2 amzn2-core 2.0 M
tuna noarch 0.13-5.amzn2.0.1 amzn2-core 139 k
tuned noarch 2.8.0-5.amzn2.0.1 amzn2-core 234 k
zlib-devel x86_64 1.2.7-17.amzn2.0.2 amzn2-core 50 k

Transaction Summary
=====
Install 12 Packages (+39 Dependent packages)

Total download size: 58 M
Installed size: 182 M
Is this ok [y/d/N]:
```



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...

polkit.x86_64 0:0.112-18.amzn2.1
polkit-pkla-compat.x86_64 0:0.1-4.amzn2.0.2
powertop.x86_64 0:2.3-12.amzn2.0.1
python-decorator.noarch 0:3.4.0-3.amzn2
python-ethtool.x86_64 0:0.8-5.amzn2.0.2
python-gobject-base.x86_64 0:3.22.0-1.amzn2.1
python-linux-procfs.noarch 0:0.4.9-3.amzn2
python-perf.x86_64 0:4.14.121-109.96.amzn2
python-pyudev.noarch 0:0.15-9.amzn2
python-schedutils.x86_64 0:0.4-6.amzn2.0.2
systemtap.x86_64 0:3.2-4.amzn2.0.2
systemtap-client.x86_64 0:3.2-4.amzn2.0.2
systemtap-devel.x86_64 0:3.2-4.amzn2.0.2
tuna.noarch 0:0.13-5.amzn2.0.1
tuned.noarch 0:2.8.0-5.amzn2.0.1
zlib-devel.x86_64 0:1.2.7-17.amzn2.0.2

Complete!
[root@ip-172-31-45-240 ~]# systemctl status tuned
● tuned.service - Dynamic System Tuning Daemon
  Loaded: loaded (/usr/lib/systemd/system/tuned.service; enabled; vendor preset
: enabled)
    Active: inactive (dead)
[root@ip-172-31-45-240 ~]#
```

# Tools for tuning

- **Let's see what the system looks like before any changes**

# Tools for tuning

- **Let's see what the system looks like before any changes**
- **We're going to see what the kernel runtime parameters are without tuned**

# Tools for tuning

- **Let's see what the system looks like before any changes**
- **We're going to see what the kernel runtime parameters are without tuned**
- **Then we'll see the base configuration**

# Tools for tuning

- **Let's see what the system looks like before any changes**
- **We're going to see what the kernel runtime parameters are without tuned**
- **Then we'll see the base configuration**
- **Then we'll try different profiles**

A screenshot of a terminal window titled "Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...". The window shows the command "sysctl -a > sysctl.before" being run, which outputs several lines of system configuration keys being read. The terminal has a light gray background and a dark gray border. A vertical scroll bar is visible on the right side of the window.

```
[root@ip-172-31-45-240 ~]# sysctl -a > sysctl.before
sysctl: reading key "net.ipv6.conf.all.stable_secret"
sysctl: reading key "net.ipv6.conf.default.stable_secret"
sysctl: reading key "net.ipv6.conf.eth0.stable_secret"
sysctl: reading key "net.ipv6.conf.lo.stable_secret"
[root@ip-172-31-45-240 ~]#
```

A screenshot of a terminal window titled 'Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...'. The window shows the command 'wc -l sysctl.before' being run, which outputs '906 sysctl.before'.

```
[root@ip-172-31-45-240 ~]# wc -l sysctl.before
906 sysctl.before
[root@ip-172-31-45-240 ~]#
```

# TXLF 2019



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...

[[root@ip-172-31-45-240 ~]# systemctl start tuned
[[root@ip-172-31-45-240 ~]# sysctl -a > sysctl.base
sysctl: reading key "net.ipv6.conf.all.stable_secret"
sysctl: reading key "net.ipv6.conf.default.stable_secret"
sysctl: reading key "net.ipv6.conf.eth0.stable_secret"
sysctl: reading key "net.ipv6.conf.lo.stable_secret"
[root@ip-172-31-45-240 ~]# ]]
```

A screenshot of a terminal window titled 'Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...'. The window shows the command 'tuned-adm active' being run, which outputs 'Current active profile: virtual-guest'.

```
[root@ip-172-31-45-240 ~]# tuned-adm active
Current active profile: virtual-guest
[root@ip-172-31-45-240 ~]#
```



# Tools for tuning

- **Diff the before and base**

# TXLF 2019



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@ec2-18-219-197-165.us-...
```

```
< kernel.ns_last_pid = 13256
---
> kernel.ns_last_pid = 13535
142c142
< kernel.random.entropy_avail = 3141
---
> kernel.random.entropy_avail = 3129
146c146
< kernel.random.uuid = 8b5a719f-a449-4998-a7ec-694b306f0c87
---
> kernel.random.uuid = 0f3c0e59-2926-48d0-8b76-fe7c9024d3eb
155c155
< kernel.sched_min_granularity_ns = 750000
---
> kernel.sched_min_granularity_ns = 10000000
163c163
< kernel.sched_wakeup_granularity_ns = 1000000
---
> kernel.sched_wakeup_granularity_ns = 15000000
871c871
< vm.dirty_ratio = 20
---
> vm.dirty_ratio = 30
902c902
< vm.swappiness = 60
---
> vm.swappiness = 30
[root@ip-172-31-45-240 ~]#
```

# Tools for tuning

- **What options do we have?**



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...

[[root@ip-172-31-45-240 ~]# tuned-adm list
Available profiles:
- atomic-guest                                - Optimize virtual guests based on the Atomic variant
- atomic-host                                   - Optimize bare metal systems running the Atomic variant
- balanced                                      - General non-specialized tuned profile
- cpu-partitioning                             - Optimize for CPU partitioning
- default                                       - Legacy default tuned profile
- desktop                                       - Optimize for the desktop use-case
- desktop-powersave                            - Optimize for the desktop use-case with power saving
- enterprise-storage                           - Legacy profile for RHEL6, for RHEL7, please use throughput-performance profile
- laptop-ac-powersave                          - Optimize for laptop with power savings
- laptop-battery-powersave                     - Optimize laptop profile with more aggressive power saving
- latency-performance                         - Optimize for deterministic performance at the cost of increased power consumption
- network-latency                               - Optimize for deterministic performance at the cost of increased power consumption, focused on low latency network performance
- network-throughput                          - Optimize for streaming network throughput, generally only necessary on older CPUs or 40G+ networks
- oracle                                       - Optimize for Oracle RDBMS
```

```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...

st of increased power consumption
- network-latency          - Optimize for deterministic performance at the co
st of increased power consumption, focused on low latency network performance
- network-throughput        - Optimize for streaming network throughput, gener
ally only necessary on older CPUs or 40G+ networks
- oracle                   - Optimize for Oracle RDBMS
- powersave                - Optimize for low power consumption
- realtime                 - Optimize for realtime workloads
- realtime-virtual-guest   - Optimize for realtime workloads running within a
KVM guest
- realtime-virtual-host    - Optimize for KVM guests running realtime worklo
ads
- sap-hana                 - Optimize for SAP
- sap-hana-vmware          - Optimize for SAP running inside a VMware guest
- sap-netweaver             - Optimize for SAP NetWeaver
- server-powersave          - Optimize for server power savings
- spindown-disk             - Optimize for power saving by spinning-down rotat
ional disks
- throughput-performance    - Broadly applicable tuning that provides excellen
t performance across a variety of common server workloads
- virtual-guest              - Optimize for running inside a virtual guest
- virtual-host               - Optimize for running KVM guests

Current active profile: virtual-guest
[root@ip-172-31-45-240 ~]#
```

# Tools for tuning

- **Let's play around a bit!**



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...  
[root@ip-172-31-45-240 ~]# tuned-adm profile balanced  
[root@ip-172-31-45-240 ~]# sysctl -a > sysctl.balanced  
sysctl: reading key "net.ipv6.conf.all.stable_secret"  
sysctl: reading key "net.ipv6.conf.default.stable_secret"  
sysctl: reading key "net.ipv6.conf.eth0.stable_secret"  
sysctl: reading key "net.ipv6.conf.lo.stable_secret"  
[root@ip-172-31-45-240 ~]#
```

# TXLF 2019



```
Key Pairs — root@ip-172-31-45-240:~ — ssh -i ThomassKeyPair.pem ec2-user@...
< fs.dentry-state = 66374      54206  45    0    0    0
---
> fs.dentry-state = 68633      56450  45    0    0    0
17,19c17,19
< fs.file-nr = 1088      0    96703
< fs.inode-nr = 54736     340
< fs.inode-state = 54736     340      0    0    0    0    0
---
> fs.file-nr = 1184      0    96703
> fs.inode-nr = 55711     340
> fs.inode-state = 55711     340      0    0    0    0    0
107c107
< kernel.ns_last_pid = 13256
---
> kernel.ns_last_pid = 14684
142c142
< kernel.random.entropy_avail = 3141
---
> kernel.random.entropy_avail = 3184
146c146
< kernel.random.uuid = 8b5a719f-a449-4998-a7ec-694b306f0c87
---
> kernel.random.uuid = 7c2e9a10-494e-4740-83c9-a66163af057e
[root@ip-172-31-45-240 ~]#
```



# Tools for tuning

- Why do you suppose nothing changed in sysctl?

# Tools for tuning

- Why do you suppose nothing changed in sysctl?
- It may have changed the I/O scheduler, CPU frequency, etc.

# Tools for tuning

- Why do you suppose nothing changed in sysctl?
- It may have changed the I/O scheduler, CPU frequency, etc.
- Since this is a cloud instance, some of the changes are not available

# Conclusion

# Conclusion

- **We've really only scratched the surface of tuning**

# Conclusion

- **We've really only scratched the surface of tuning**
- **We can tune networking, file services, and so on**

# Conclusion

- **We've really only scratched the surface of tuning**
- **We can tune networking, file services, and so on**
- **But you should have a basic understanding of what and how you can inspect and change things on Linux**



# THANK YOU!



# Questions?