
Site Reliability Engineering with Prometheus

Eddy Reyes • eddy@mindsight.io

SRE vs Devops

Devops

- Set of Practices and Principles for Better Cloud Applications
- Broad, Transformational, Managerial

SRE

- Narrower Focus on Reliability
 - A Way to Implement Devops, with Specific Recommendations
-

Reliability At Scale


Metrics

- Rates
- Errors
- Distributions / Durations

Process

- Service Level Objectives
 - Blameless Postmortems
-

100% ... Good Luck!



Not Really Attainable

- Your dependencies are <100%
 - Never change anything again
 - Prohibitive cost for little additional value
-

Service Level Objectives (SLOs)

- Percentage (<100%)
- Over Period of Time
 - e.g. “30 days”

Based on Acceptable Quality to
Your **Customers**

SLAs vs. SLOs

SLA

- Contractual obligation of availability or refund



SLO

- Internal metric, based on observation
- Can change / be iterated upon



SLO Benefits

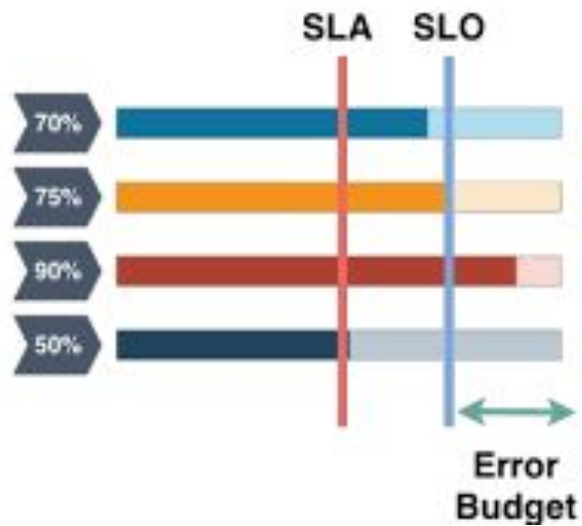
Org Interop

- Business vs. Engineering Criteria for Success
- Customer-focused and business-driven


Engineering Enabler

- <100% -> Acceptable Level of Failure
 - Learning / Iteration
 - SLO Failure Is Now the Only Alerting That Matters
 - Alert Fatigue
-

Error Budgets



Latency Error Budgets



How To Calculate

- Choose Maximum Latency X
- Error budget:
 - $\#Samples * 1-SLO < \#Samples(>X)$

Error Budget Benefits

Process Safety

- Is it safe to deploy right now?
- Does an alert need to wake me up at 3am?

Experimentation

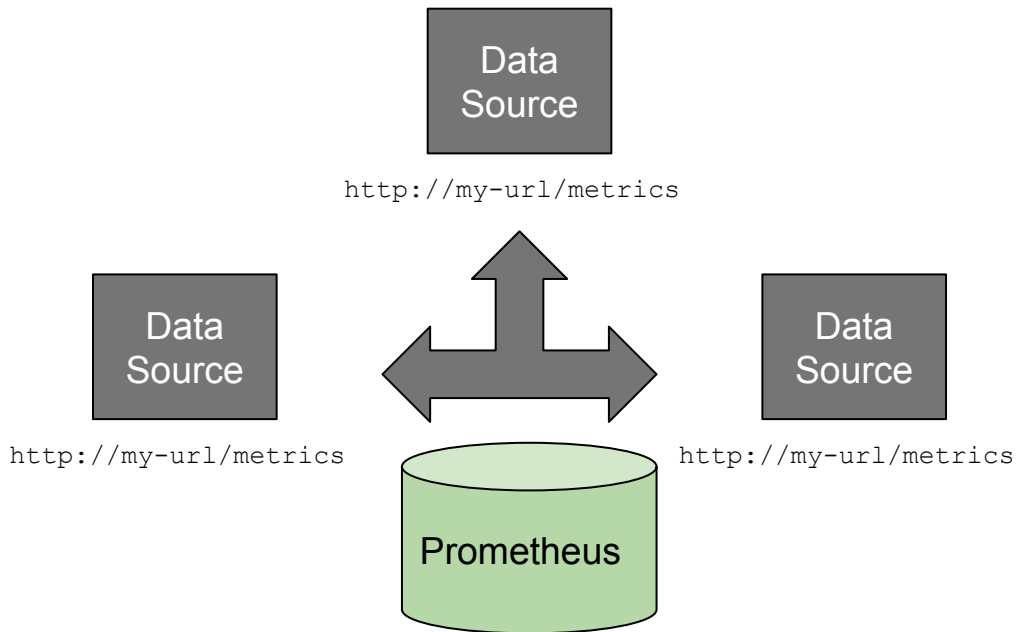
- Chaos Engineering
- Example: Facebook
 - Intentional crashing for UX research

Where Does the Data Come From?

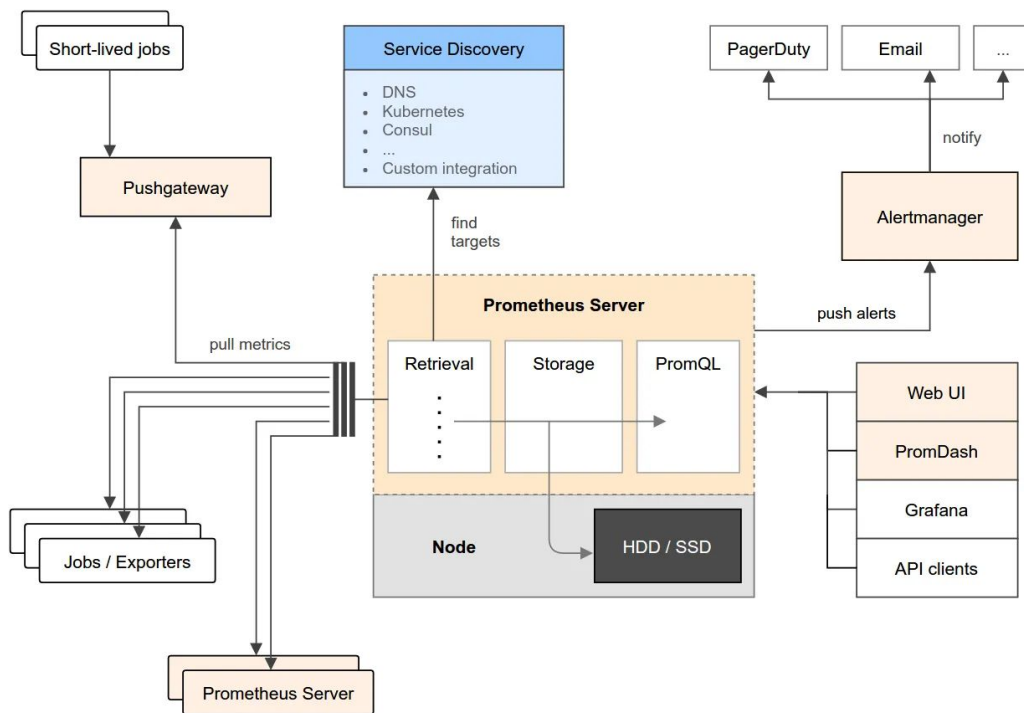
Service Level Indicators (SLIs)

- Telemetry Used To Compute SLOs
- Error Rates:
 - Success / total (over SLO period)
- Latencies
 - 95% of all latencies below X

Prometheus Pulls Metrics



Prometheus



Prometheus Data Model

Queries - PromQL

```
series_name{tag1="A"}
```

Result (Instant Vector):

```
series_name{tag1="A" job="mytest"} 10.81  
series_name{tag1="A" job="othertest"} 100.17
```

Prometheus Data Model

Data Types

- Counters
 - Monotonically increasing
 - e.g. request counters
 - Gauges
 - Instantaneous reading
 - e.g. CPU usage %
 - Histograms / Samples
 - Values kept as distributions
 - e.g. Request latency
-

SLIs With Prometheus

Error Rate

```
sum(rate(slo_errors[10m])) by (job) / sum(rate(slo_requests[10m]))  
by (job)
```


Latency Quantile

```
histogram_quantile(0.9,  
sum(rate(http_request_duration_seconds_bucket[10m])) by (job, le))
```

SRE-Friendly Alerts


- Given Error Budget, Compute the Burn Rates
- Am I About to Fail SLO In:
 - 1 hour? WAKE UP!!
 - 1 day? WAKE UP!!
 - 1 week? Look on Monday...

Alert On Burn Rate



Burn rate	Error rate for a 99.9% SLO	Time to exhaustion
1	0.1%	30 days
2	0.2%	15 days
10	1%	3 days
1,000	100%	43 minutes

Alert Sensitivity



```
expr: (
  job:slo_errors_per_request:ratio_rate1h{job="myjob"} > (14.4*0.001)
  and
  job:slo_errors_per_request:ratio_rate5m{job="myjob"} > (14.4*0.001)
)
or
(
  job:slo_errors_per_request:ratio_rate6h{job="myjob"} > (6*0.001)
  and
  job:slo_errors_per_request:ratio_rate30m{job="myjob"} > (6*0.001)
)
severity: page

expr: (
  job:slo_errors_per_request:ratio_rate24h{job="myjob"} > (3*0.001)
  and
  job:slo_errors_per_request:ratio_rate2h{job="myjob"} > (3*0.001)
)
or
(
  job:slo_errors_per_request:ratio_rate3d{job="myjob"} > 0.001
  and
  job:slo_errors_per_request:ratio_rate6h{job="myjob"} > 0.001
)
severity: ticket
```

Mindsight Simplifies SRE

The screenshot displays the Mindsight SRE dashboard interface. A modal titled "Add New Application" is open, showing a progress bar with four steps: "Add Application" (checked), "Define SLO" (checked), "Define SLIs (Data Sources)" (checked), and "Configure Alerts" (active). The modal contains three identical configuration sections, each with an "Error Budget Burn %" slider and a "Within Time Window" section with "Days" and "Hours" radio buttons. The background dashboard shows a "Needs Attention" section with "SLOs At Risk" and an "Applications Overview" section with a table of applications.

App Name	Since
Public API	08/20/2019
Backend	08/20/2019

Instrumenting Infrastructure Into Prometheus

- Many tools help here
- Promising telemetry technology:
 - Kubernetes + Istio
 - Available on Google Cloud

Thank You!!

Eddy Reyes

eddy@mindsight.io

@0x000edd1e
